

Functions of operating system

- File management
- Device management
- Memory management
- Process management

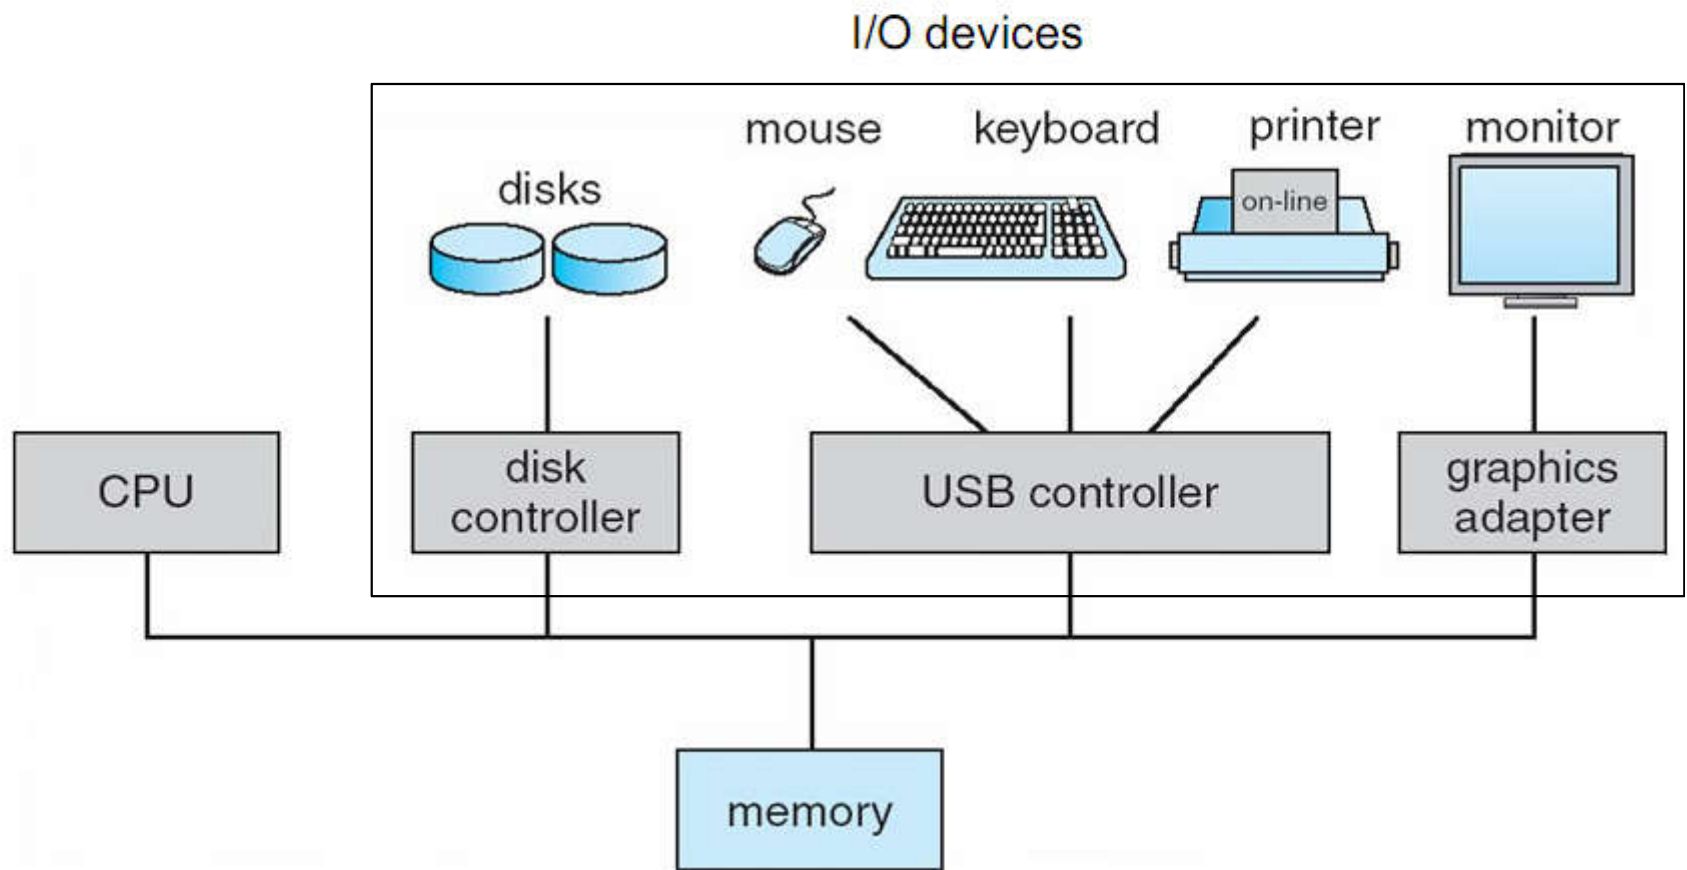
Functions of operating system

- 文件管理
 - **File management system:** FAT32 NTFS EXT3 EXT4 ZFS Btrfs
 - 控制文件的访问
 - 管理文件的创建、删除和修改
 - 给文件命名
 - 管理文件的存储

Functions of operating system


- 设备管理
 - 访问输入输出设备
 - 输入输出设备存在着数量和速度限制
 - 提高设备使用效率

Computer System Organization



Speeding up I/O: Direct Memory Access (DMA)

- Data moved directly between I/O devices and memory
- CPU can work on other tasks



Functions of operating system

- 内存管理

- 单道程序

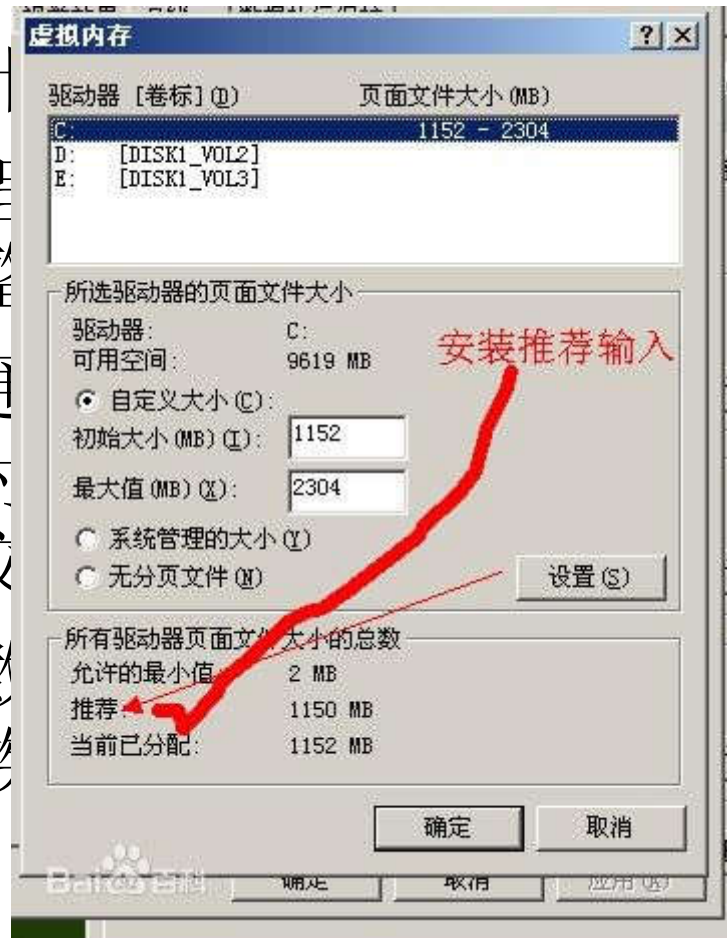
- 程序需全载入内存
 - 一个程序运行时，其他程序不能运行

- 多道程序

- 同一时刻可载入多个程序
 - 模式：非交换、交换
 - 非交换：程序在运行期间始终驻留在内存
 - 交换：运行过程中，程序可以在内存和硬盘间多次交换数据

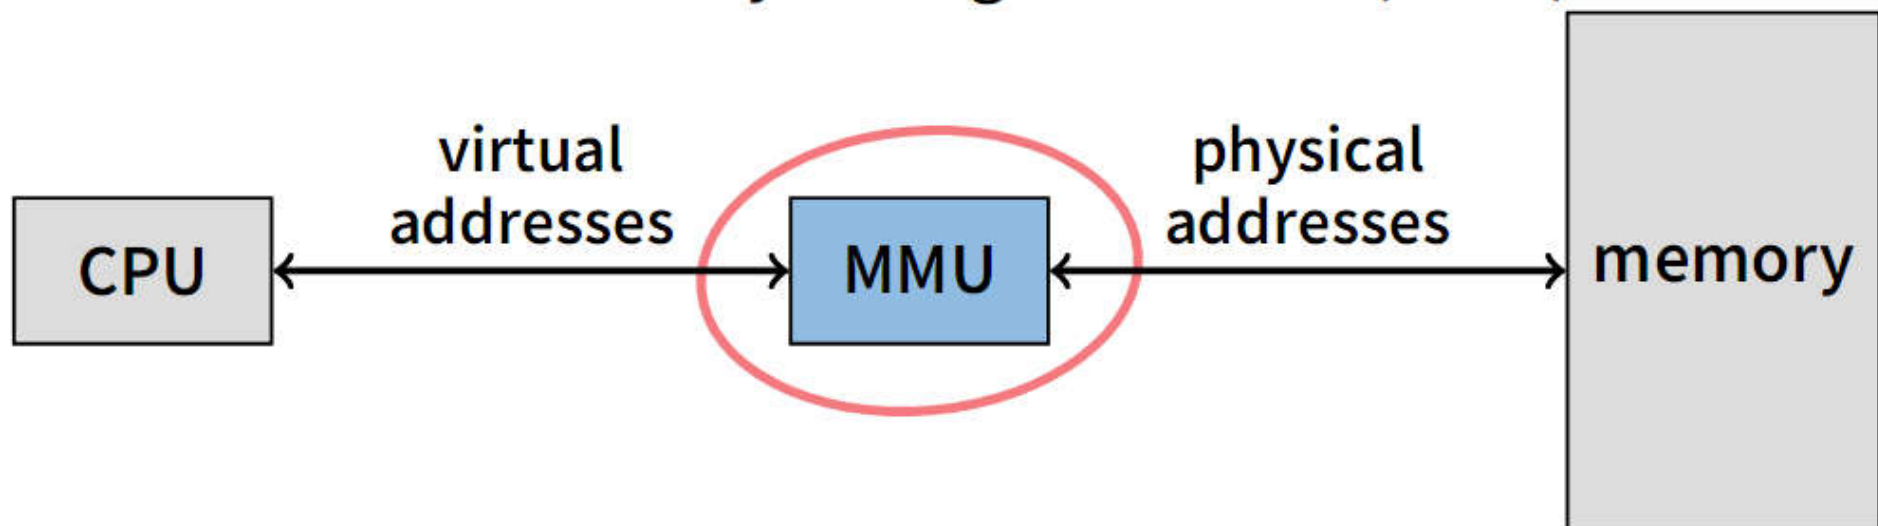
- 虚拟内存

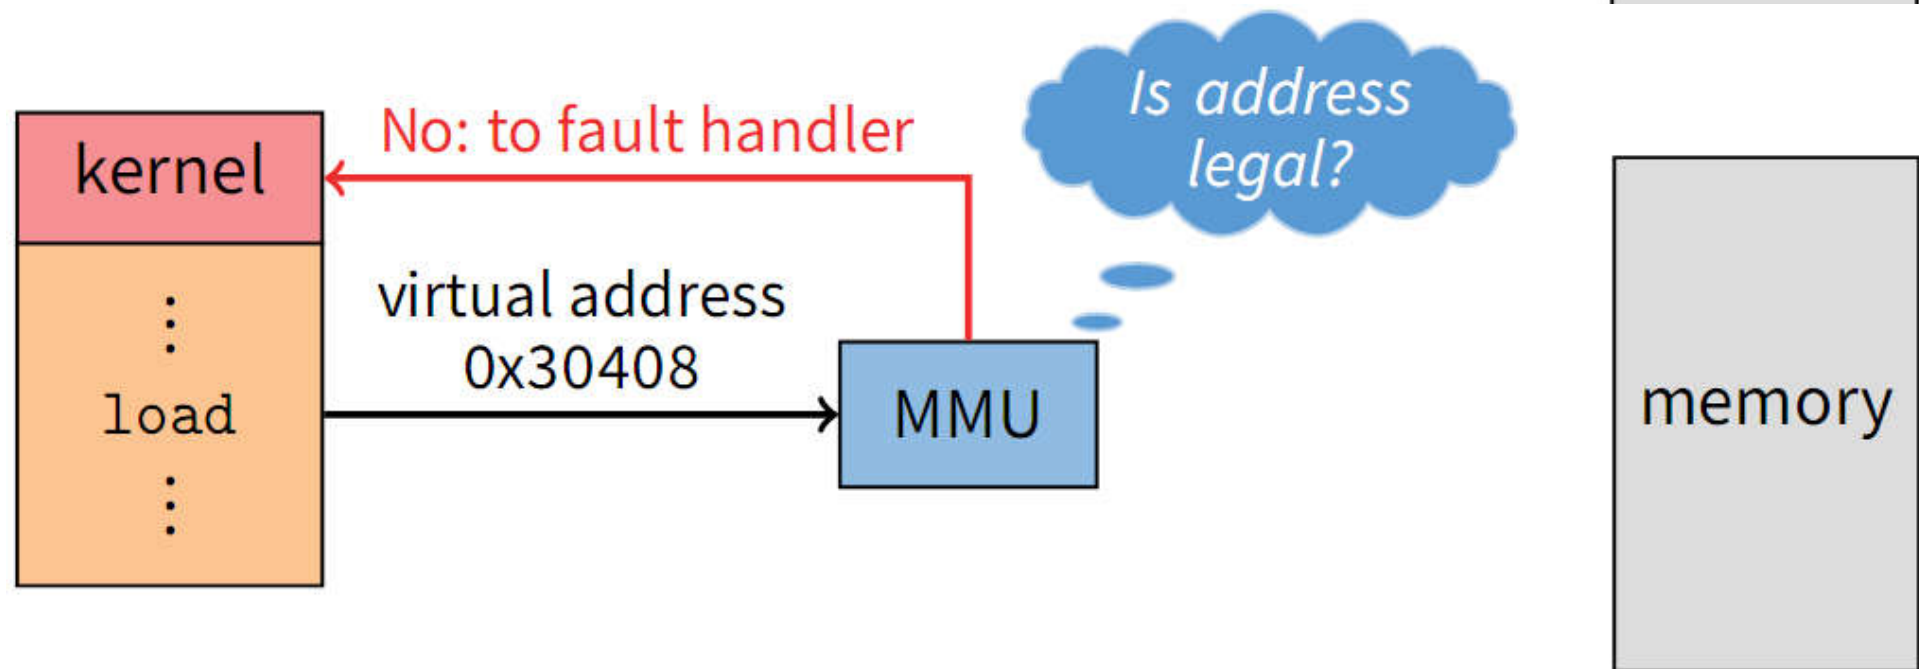
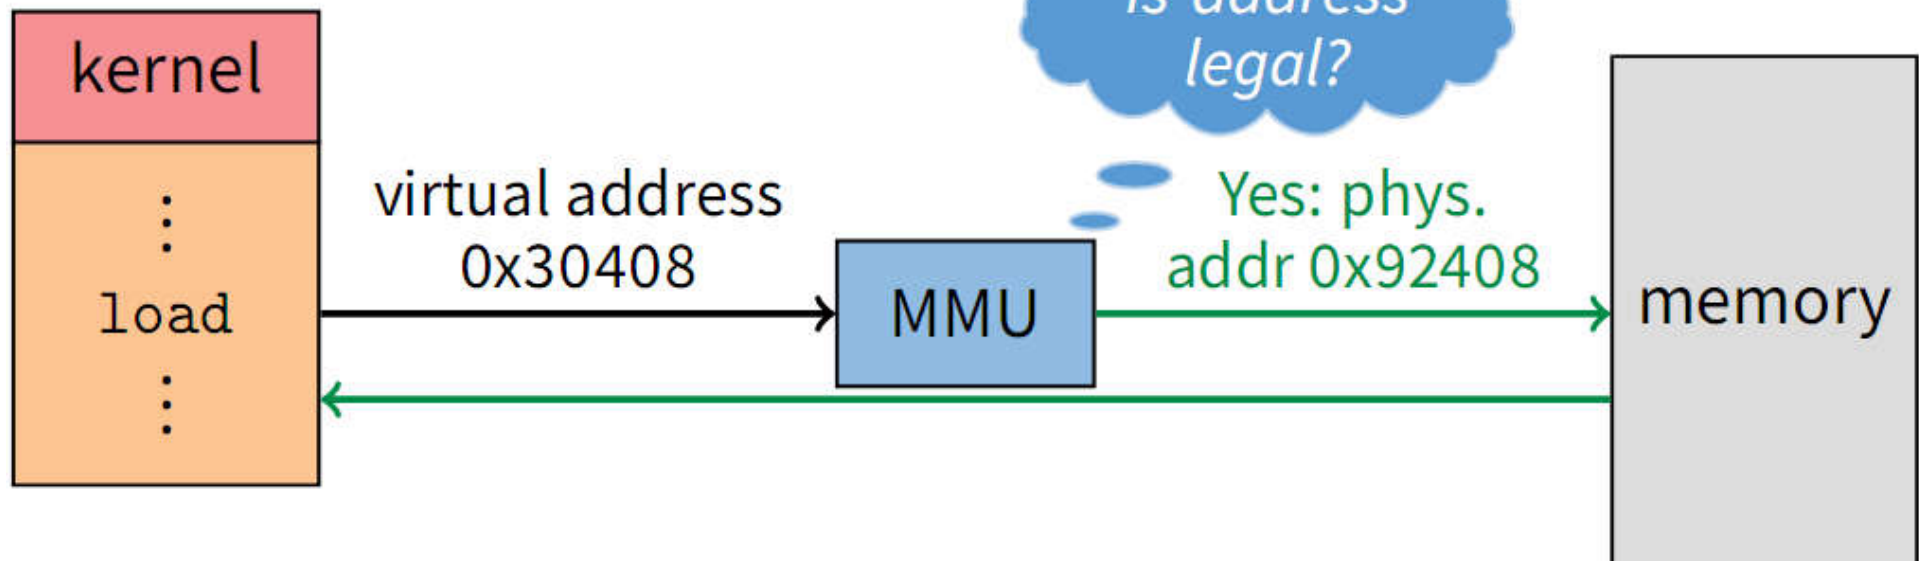
- 虚拟内存是计
- 它使得应用程序一个连续完整
- 实际上，它通
还有部分暂时
时进行数据交
- 目前，大多数
Windows家族
间”等



一种技术
可用的内存（
理内存碎片，
器上，在需要
以内存，如
ux的“交换空

- Programs load/store to **virtual addresses**
- Actual memory uses **physical addresses**
- VM Hardware is Memory Management Unit (**MMU**)





Process进程

- What is a process

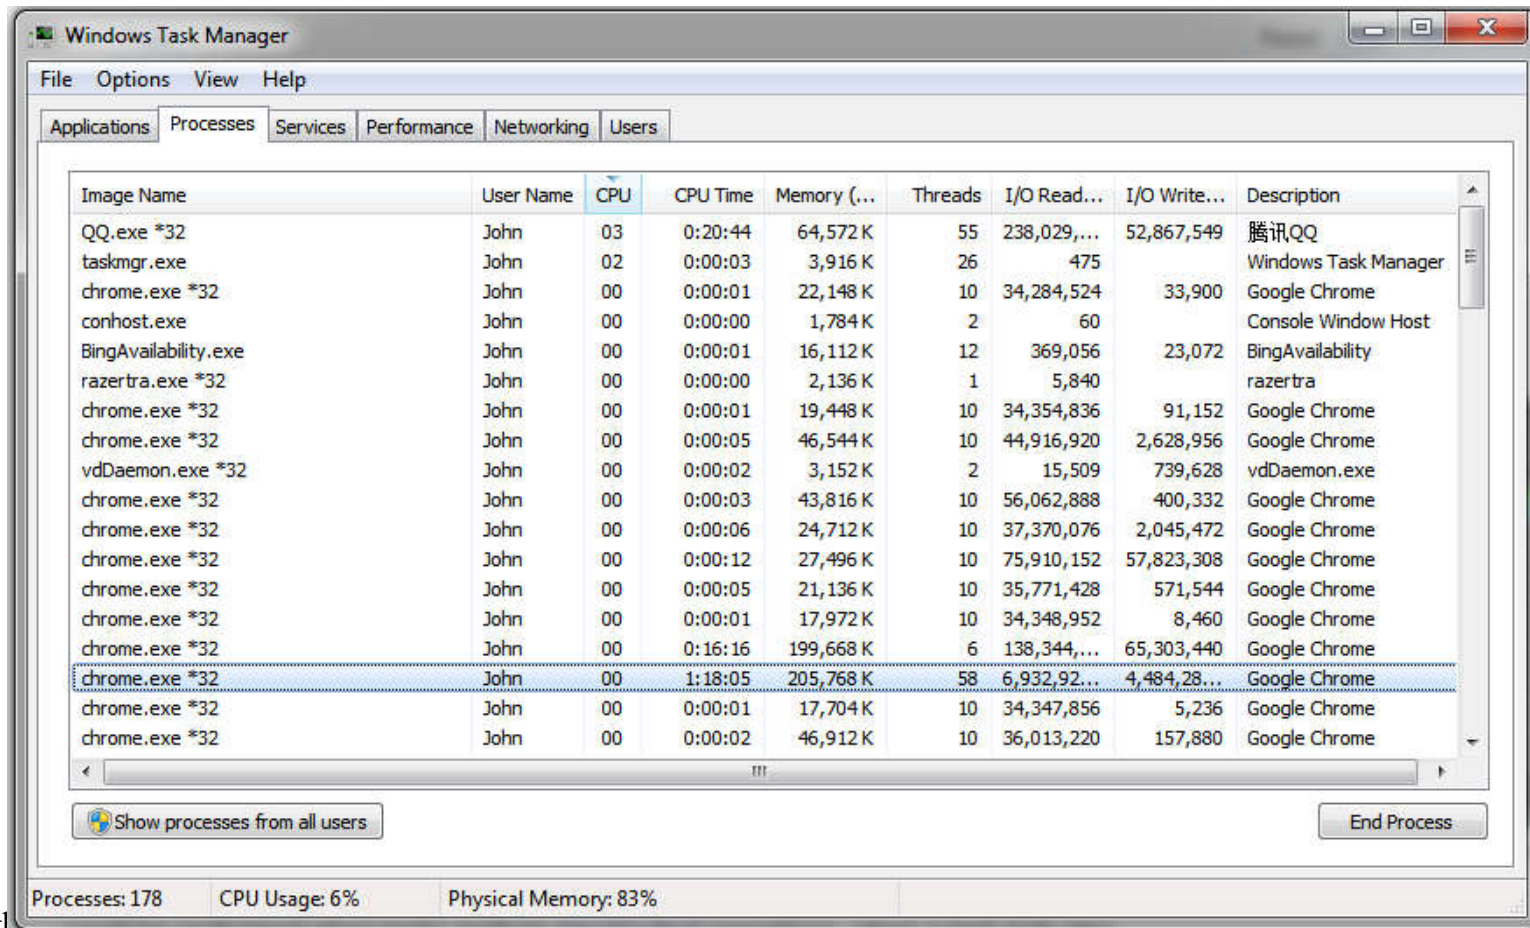
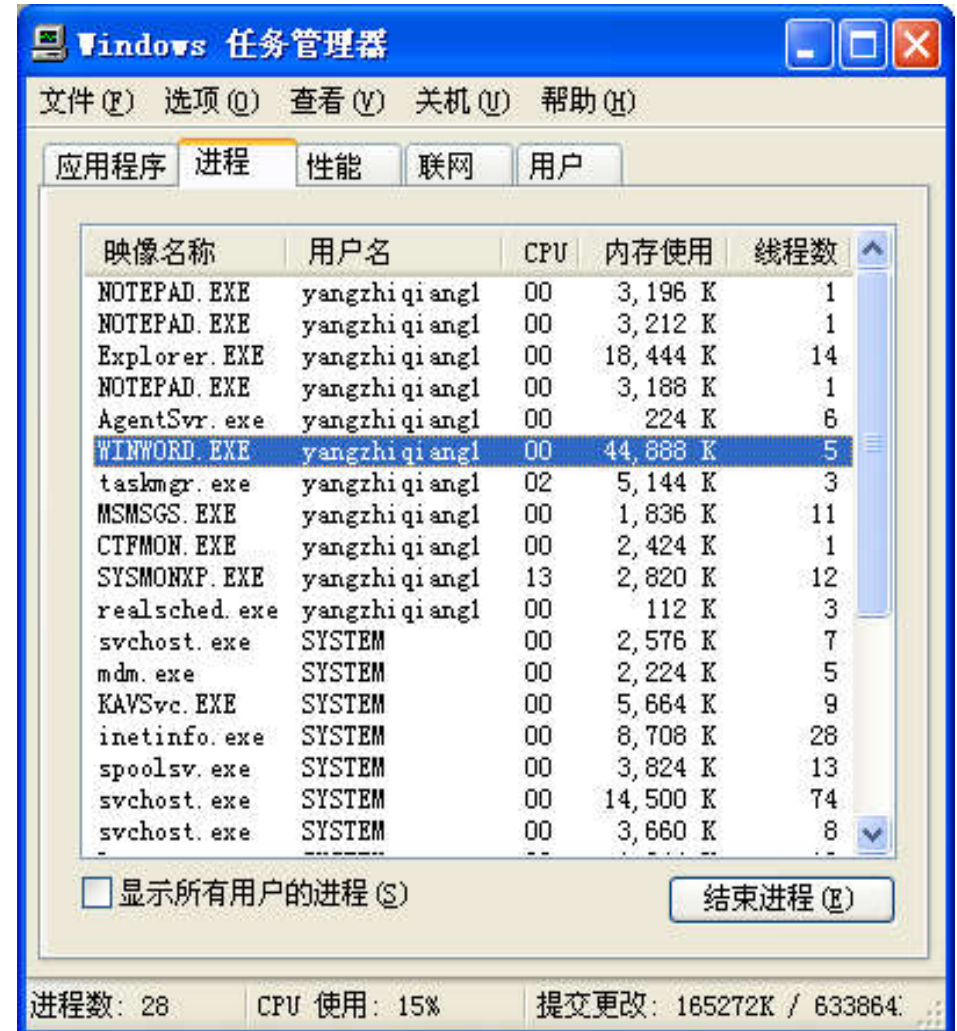
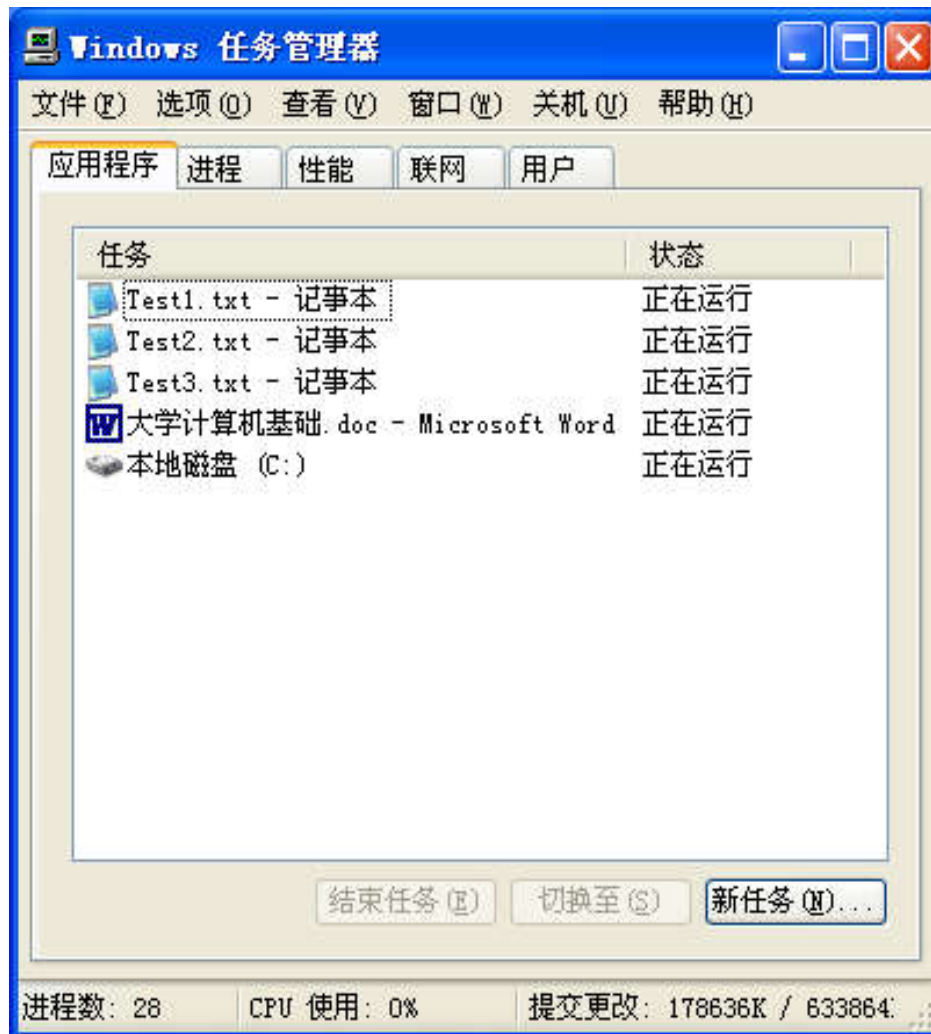


Image Name	User Name	CPU	CPU Time	Memory (...)	Threads	I/O Read...	I/O Write...	Description
QQ.exe *32	John	03	0:20:44	64,572 K	55	238,029,...	52,867,549	腾讯QQ
taskmgr.exe	John	02	0:00:03	3,916 K	26	475		Windows Task Manager
chrome.exe *32	John	00	0:00:01	22,148 K	10	34,284,524	33,900	Google Chrome
conhost.exe	John	00	0:00:00	1,784 K	2	60		Console Window Host
BingAvailability.exe	John	00	0:00:01	16,112 K	12	369,056	23,072	BingAvailability
razertra.exe *32	John	00	0:00:00	2,136 K	1	5,840		razertra
chrome.exe *32	John	00	0:00:01	19,448 K	10	34,354,836	91,152	Google Chrome
chrome.exe *32	John	00	0:00:05	46,544 K	10	44,916,920	2,628,956	Google Chrome
vdDaemon.exe *32	John	00	0:00:02	3,152 K	2	15,509	739,628	vdDaemon.exe
chrome.exe *32	John	00	0:00:03	43,816 K	10	56,062,888	400,332	Google Chrome
chrome.exe *32	John	00	0:00:06	24,712 K	10	37,370,076	2,045,472	Google Chrome
chrome.exe *32	John	00	0:00:12	27,496 K	10	75,910,152	57,823,308	Google Chrome
chrome.exe *32	John	00	0:00:05	21,136 K	10	35,771,428	571,544	Google Chrome
chrome.exe *32	John	00	0:00:01	17,972 K	10	34,348,952	8,460	Google Chrome
chrome.exe *32	John	00	0:16:16	199,668 K	6	138,344,...	65,303,440	Google Chrome
chrome.exe *32	John	00	1:18:05	205,768 K	58	6,932,92...	4,484,28...	Google Chrome
chrome.exe *32	John	00	0:00:01	17,704 K	10	34,347,856	5,236	Google Chrome
chrome.exe *32	John	00	0:00:02	46,912 K	10	36,013,220	157,880	Google Chrome

Processes: 178 CPU Usage: 6% Physical Memory: 83%

Processes



Process


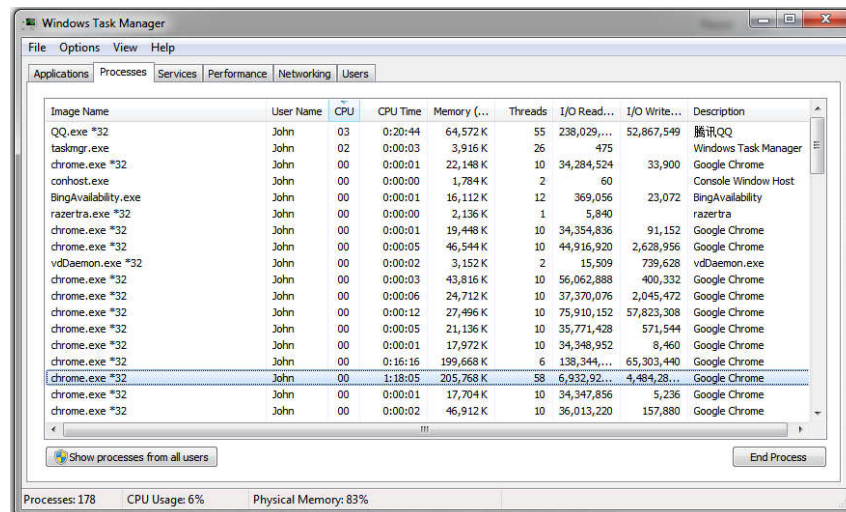
- Each process represents a task
- Jobs vs Tasks
 - A collection of tasks that is used to perform a computation is known as a *job* (MSDN)
- Is a **process** the same as a **program**?

A Program vs. a Process

- **Program**: a set of instructions, e.g., notepad.c, notepad.exe
- **Process**: activity of executing a program
- A program can be run multiple times, each instance/activity called a **process**
- **Process State**: Current status of the activity
 - Program counter
 - General purpose registers
 - Related portion of main memory

Process table

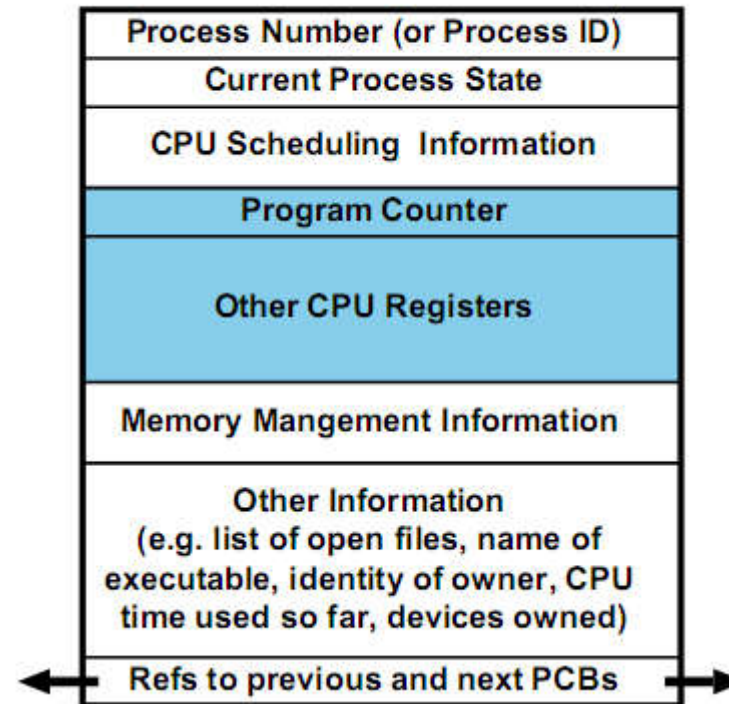
- A table of all the processes maintained by the operating system.
- Each entry is a process and its descriptions (PCB).



Process control block (PCB)

- "the manifestation of a process in an operating system"
 - Process identification
 - Processor state data (the status of a process, saved registers, program counter etc.)
 - Process control data (scheduling state, privileges etc.)

Process Control Block



The secret of concurrent execution

- What do you have to do when switching from one ongoing task to another?
 - Simply stop the current task and turn to another
 - Record the status of the current task and suspend it and turn to another
- Context switch


What is a context?

- Snapshot of current status of a process (PCB)
 - A process identifier, or PID
 - Register values, Program Counter value
 - The memory space, I/O, files for the process
 - Can be saved and resumed as if the process is not interrupted
- Another meaning: execution *state* of the process
 - Ready: ready for execution
 - Waiting: waiting for some I/O
 - Complete: finished process

Context switch

- The process of storing and restoring the state (context) of a process so that execution can be resumed from the same point at a later time.
- This enables multiple processes to share a single CPU and is an essential feature of a multitasking operating system.

Figure 3.6 Time-sharing between process A and process B



Who is responsible for context switching?

- **Process management**
 - **Scheduler (调度):** Adds new processes to the process table and removes completed processes from the process table
 - **Dispatcher (分派):** Controls the allocation of time slices to the processes in the process table

When to switch?

- **Interrupt** a signal to the processor emitted by hardware or software indicating an event that needs immediate attention.
- The *interrupt handler* (part of dispatcher) starts after the interrupt to perform context switch
- Modern architectures are interrupt driven.
 - Software interrupt (I/O)
 - Hardware interrupt (press a key)

Scheduler

- Determines which processes should be considered for execution based on some priorities or concerns
 - Using process table for administration
- Process table
 - Ready or waiting
 - Priority
 - Non-scheduling information: memory pages, etc.

Dispatcher

- Gives time slices to a process that is ready
- Executes a **context switch** when the running process's time slice is over
 - *Time slice*: a time segment for each execution
 - *Interrupt*: the signal generated by a hardware timer to indicate the end of a time slice.
 - The *interrupt handler* (part of dispatcher) starts after the interrupt to perform context switch

Context Switch (process switch)

1. Get an interrupt from timer
2. Go to the interrupt handler
 - a. Save the context of process A
 - b. Find a process ready to run (Assume that is process B)
 - c. Load the context of process B
3. Start (continue) process B


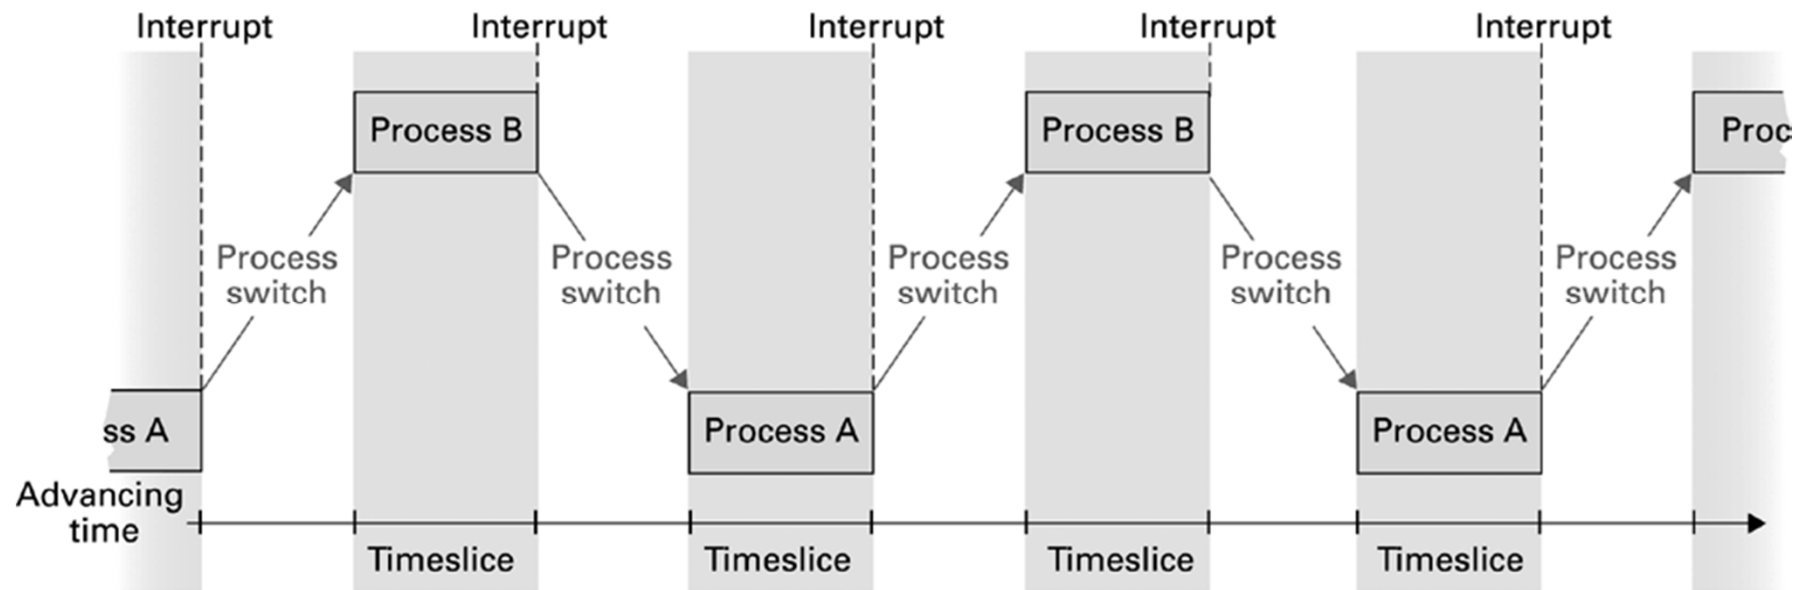



Figure 3.6 Time-sharing between process A and process B



Thread

- A task existing within a process that allows multiple independent instances to be executed concurrently
 - Multiple threads share resources such as memory, program code, ...
 - Each thread has its own program counter, registers, and stack (local memory)
- The context switch of threads is much faster than that of processes



Windows 任务管理器

文件(F) 选项(O) 查看(V) 关机(U) 帮助(H)

应用程序 进程 性能 联网 用户

映像名称	用户名	CPU	内存使用	线程数
NOTEPAD.EXE	yangzhiqiang	00	3,196 K	1
NOTEPAD.EXE	yangzhiqiang	00	3,212 K	1
Explorer.EXE	yangzhiqiang	00	18,444 K	14
NOTEPAD.EXE	yangzhiqiang	00	3,188 K	1
AgentSvr.exe	yangzhiqiang	00	224 K	6
WINWORD.EXE	yangzhiqiang	00	44,888 K	5
taskmgr.exe	yangzhiqiang	02	5,144 K	3
MSMSG.S.EXE	yangzhiqiang	00	1,836 K	11
CTFMON.EXE	yangzhiqiang	00	2,424 K	1
SYSMONXP.EXE	yangzhiqiang	13	2,820 K	12
realsched.exe	yangzhiqiang	00	112 K	3
svchost.exe	SYSTEM	00	2,576 K	7
mdm.exe	SYSTEM	00	2,224 K	5
KAVSvc.EXE	SYSTEM	00	5,664 K	9
inetinfo.exe	SYSTEM	00	8,708 K	28
spoolsv.exe	SYSTEM	00	3,824 K	13
svchost.exe	SYSTEM	00	14,500 K	74
svchost.exe	SYSTEM	00	3,660 K	8

☐ 显示所有用户的进程(S)

结束进程(E)

进程数: 28 CPU 使用: 15% 提交更改: 165272K / 633864

查看(V) 帮助(H)

立即刷新(R)

更新速度(U)

选择列(S)...

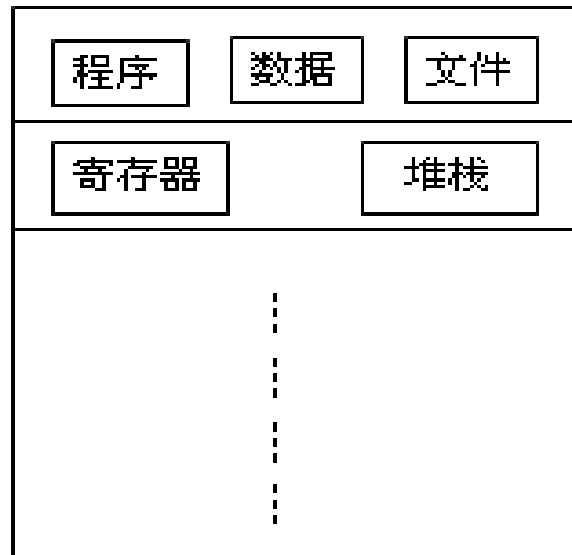
选择列

请选择“任务管理器”进程页上将显示的列。

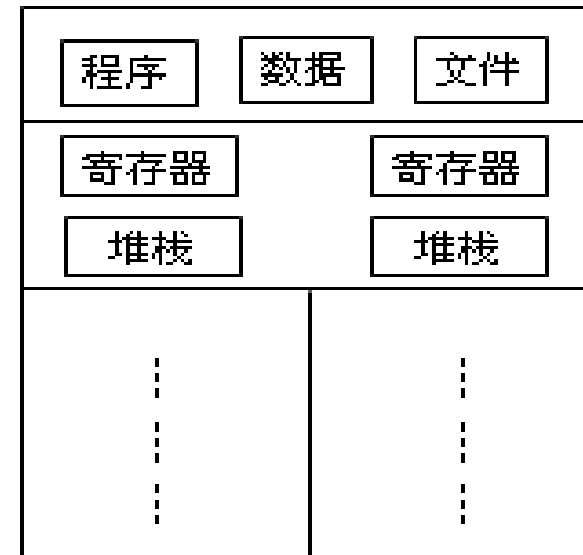
<input checked="" type="checkbox"/> 映像名称(I)	<input type="checkbox"/> 页面错误增量(A)
<input type="checkbox"/> PID (进程标识符)(P)	<input type="checkbox"/> 虚拟内存大小(V)
<input checked="" type="checkbox"/> CPU 使用(C)	<input type="checkbox"/> 页面缓冲池(G)
<input type="checkbox"/> CPU 时间(E)	<input type="checkbox"/> 非页面缓冲池(Q)
<input checked="" type="checkbox"/> 内存使用(M)	<input type="checkbox"/> 基本优先级(R)
<input type="checkbox"/> 内存使用增量(U)	<input type="checkbox"/> 句柄计数(H)
<input type="checkbox"/> 内存使用高峰值(K)	<input checked="" type="checkbox"/> 线程计数(T)
<input type="checkbox"/> 页面错误(F)	<input type="checkbox"/> GDI 对象(O)
<input type="checkbox"/> USER 对象(U)	<input type="checkbox"/> I/O 写入
<input type="checkbox"/> I/O 读取	<input type="checkbox"/> I/O 写入字节
<input type="checkbox"/> I/O 读取字节	<input type="checkbox"/> I/O 其他
<input checked="" type="checkbox"/> 会话 ID(S)	<input type="checkbox"/> I/O 其他字节
<input checked="" type="checkbox"/> 用户名(N)	

确定 取消

A Thread vs. a Process



单线程



多线程

Share more

High Concurrency

Fast Switch

Exercises

- Suppose an OS allocates time slices in 10 millisecond units and the time required for a context switch is negligible. How many processes can obtain a time slice in one second?

Competition for Resources

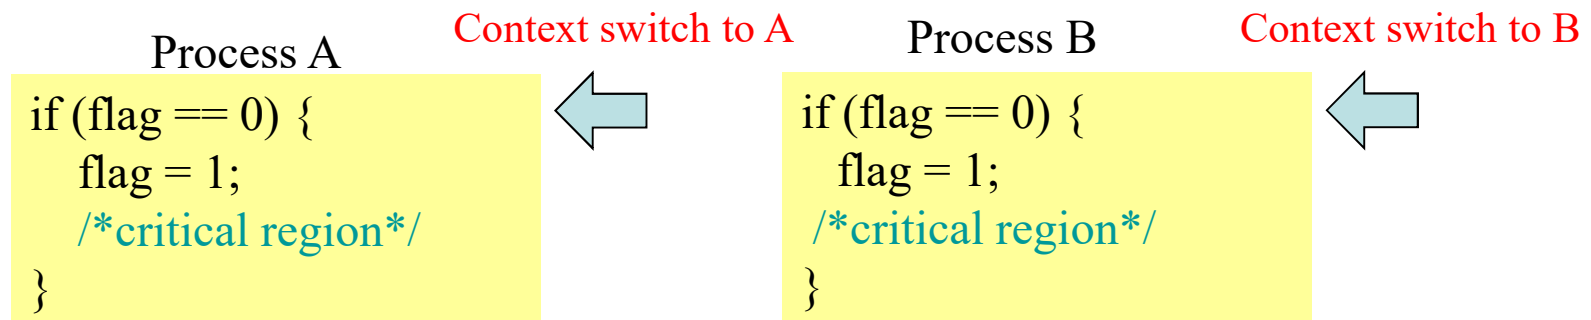
- What are *resources*?
 - CPU, memory, files, peripheral devices, ...
- In a multitasking system, resources are shared by processes
 - Some resources should not be employed by more than one process simultaneously
 - E.g., printer

Handling Competition for Resources

- **Semaphore:** A “control flag”
- **Critical Region:** A group of instructions that should be executed by only one process at a time
- **Mutual exclusion:** Requirement for proper implementation of a critical region

First Algorithm

- Use a flag (a global memory address)
 - flag=1: the critical region is occupied
 - flag=0: no process is in the critical region
- Problem:




- Both processes get into the critical region

Solutions

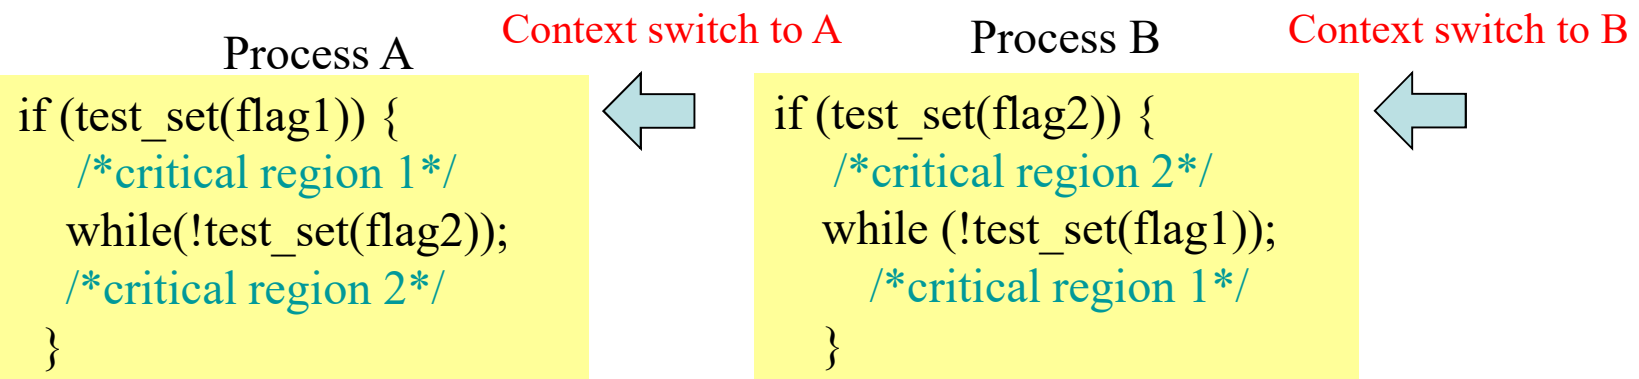
- Testing&setting the flag must be completed without interruption
1. Use `disable_Interrupt()` to prevent context switch during the flag test and set process.

```
Disable_Interrupt();
if (flag == 0) {
    flag = 1;
    Enable_Interrupt();
    /*critical region*/
}
Enable_Interrupt();
```
 2. A machine instruction called “test-and-set” which cannot be interrupted
- Semaphore: a properly implemented flag

Another Problem: Deadlock




- Example:
 - A is in critical region 1, and waits to enter critical region 2
 - B is in critical region 2, and waits to enter critical region 1



Deadlock

- Processes block each other from continuing
- Conditions required for deadlock
 1. Competition for non-sharable resources
 2. Resources requested on a partial basis
 3. An allocated resource can not be forcibly retrieved

Remove any one of the conditions
can resolve the deadlock.



Solutions

Which condition is removed?

1. Kill one of the process
2. Processes need to request all the required resources at one time
3. Spooling
 - For example, stores the data to be printed and waits the printer available
4. Divide a file into pieces so that it can be altered by different processes

Exercises

- There is a bridge that only allows one car to pass. When two cars meet in the middle, it causes “deadlock”. The following solutions remove which conditions
 - Do not let a car onto the bridge until the bridge is empty.
 - If cars meet, make one of them back up.
 - Add a second lane to the bridge.
- What’s the drawback of solution 1?

Virtual machine虚拟机

Physical Machine

